

Optimizing Early Detection of Plant Diseases: A Web Application Using MobileNet as a CNN Network

Ismail Jabri², Aziza El Ouaazizi^{1,2}, Zine Eddine Lourig¹, Siham El Ghazi², Chaimae Khattabi², and Mohamed Ettahairy²

¹Laboratory of Artificial Intelligence, Data Sciences and Emergent Systems (LIASSE), National School of Engineers (ENSA), Sidi Mohamed Ben Abdellah University, Fez, Morocco

²Laboratory of Engineering Sciences (LSI), Polydisciplinary Faculty of Taza, Sidi Mohamed Ben Abdellah University, Fez, Morocco

Abstract

Plant diseases remain a persistent challenge to global agriculture, especially in regions with limited access to diagnostic infrastructure and expert support. While deep learning has shown promise in image-based classification tasks, its real-world application in farming remains underdeveloped. This study presents a lightweight, web-based deep learning model for early plant disease detection using RGB images of affected leaves. The proposed model is built upon the MobileNetV2 architecture and fine-tuned using transfer learning to balance accuracy and computational efficiency. Crucially, we sourced and curated a robust, diverse dataset from open-access agricultural repositories and academic archives. This dataset comprises 54 disease classes across 26 plant species, chosen to reflect real-world conditions. This realistic approach is essential for training the model to perform accurately on new, unseen images, ensuring it is a reliable and effective tool for farmers in the field. Data augmentation techniques were applied to further enhance model robustness. We also deployed this model in an intuitive web application that allows to users to upload leaf images, receive instant diagnostic results, and view tailored treatment suggestions. This work demonstrates that MobileNetV2, coupled with a carefully designed dataset and accessible interface, can bridge the gap between state-of-the-art AI and field-level agricultural needs, providing a scalable, user-friendly tool for real-time crop disease management and contributing to food security through technological innovation.

Keywords: Early diagnosis; Convolutional neural network; Plant diseases; MobileNet; Deep Learning

1 Introduction

Agriculture is a vital sector in many developing countries, providing food security and livelihoods. According to the FAO, plant diseases cause 20–40% of global crop losses annually, disproportionately affecting smallholder farmers lacking timely diagnostic tools [1]. Traditional diagnostic

*ismail.jabri@usmba.ac.ma

[†]aziza.elouaaazizi@usmba.ac.ma

[‡]zineeddine.louriga@usmba.ac.ma

[§]siham.elghazi@usmba.ac.ma

[¶]chaimae.khattabi@usmba.ac.ma

^{||}mohamed.ettahairy@usmba.ac.ma

methods—such as expert visual inspection, PCR, and ELISA—are accurate but often slow, costly, and dependent on specialized equipment [2–4].

Recent advances in digital agriculture, including remote sensing and AI, have enabled more efficient crop monitoring and plant disease detection. Deep learning, in particular, has shown strong performance in image-based classification tasks [5, 6]. However, most CNN models are trained on controlled datasets with consistent lighting and backgrounds [7], making them less effective in real-world scenarios. Additionally, high-performance models like ResNet and Inception are computationally intensive and not ideal for mobile use [8, 9].

To address these issues, lightweight architectures such as MobileNet have been proposed for deployment on resource-limited devices, maintaining high accuracy with lower complexity [8]. Studies using MobileNet for plant disease detection have achieved over 95% accuracy [16, 17], and some integrate it with hybrid techniques like Squeeze-and-Excitation blocks for improved results [18, 19].

Contributions : The main contributions of this study are:

1. **Real-World, Large-Scale Dataset.** We curate and balance 54 disease classes across 26 plant species—sourced from open repositories and field images—to faithfully represent real agricultural conditions.
2. **Fine-Tuned MobileNetV2 for Edge Deployment.** We adapt MobileNetV2 via transfer learning, achieving 97.40% accuracy with only 3.5 M parameters to enable efficient inference on resource-constrained devices.
3. **Browser-Accessible Web Application.** We deliver a Flask-based interface allowing users to upload leaf images, receive instant diagnostic results, and view tailored treatment suggestions—bridging the gap between AI research and field deployment.

2 Related Work

The early and accurate detection of plant diseases is critical to reducing crop losses, maintaining food security, and enabling precision agriculture. With the proliferation of imaging devices and recent advances in artificial intelligence, particularly deep learning, numerous researchers have explored automated plant disease identification through image classification. Among deep learning approaches, Convolutional Neural Networks (CNNs) have shown particular promise due to their ability to model complex visual features, making them especially well-suited for leaf-based disease diagnosis.

Kulkarni [10] was among the first to propose a Deep Convolutional Neural Network (DCNN)-based approach for plant disease recognition, using a small custom dataset with five crops and three disease classes per crop [11]. The model achieved an accuracy of 93.5%, establishing the potential of CNNs in this domain. However, the dataset lacked diversity, and the system was not evaluated under real-world field conditions, which are often affected by variations in lighting, backgrounds, and leaf occlusion.

A major contribution to this field came with the introduction of the PlantVillage dataset by Hughes and Salathé [12], which provided over 50,000 labeled images of healthy and diseased leaves across several crops. This dataset enabled the training and evaluation of more robust models and established a common benchmark for researchers. Mohanty et al. [13] trained AlexNet and GoogLeNet on PlantVillage and achieved up to 99.35% classification accuracy across 14 crop species and 26 diseases. Despite the high performance, the dataset consists mainly of laboratory-acquired images with uniform backgrounds, reducing the generalizability of models to real-world scenarios.

Ferentinos [7] extended this line of research by comparing the performance of AlexNet, VGG, and ResNet across 58 plant disease classes. The VGG-based model reported the highest accuracy of 99.53%, reinforcing the effectiveness of deep CNN architectures. Nonetheless, the study was similarly constrained by dataset homogeneity and did not address resource limitations or deployment mechanisms on edge devices or in low-connectivity environments.

In efforts to make plant disease detection systems more accessible to end-users such as farmers, some studies have integrated CNNs into mobile applications [24]. Mique Jr. and Eusebio [14] developed a mobile tool for identifying rice insect pests and diseases using a CNN, achieving 91.8% accuracy. However, the app was built for a single crop and lacked scalability or field validation. Ramcharan et al. [15] took a more practical approach by building and field-testing a deep learning model for cassava disease detection in Tanzania. The model, trained on real user-captured images, achieved over 95% accuracy and proved effective in real-world settings [8]. Nevertheless, this work focused on a single platform (mobile), lacked cross-platform support, and did not explore computational trade-offs or web deployment [9].

3 Methodology

3.1 Data Acquisition

To build an accurate plant disease classifier, we created a dataset with images of 26 plant species and 54 disease classes, mainly captured using RGB cameras showing symptomatic leaves on plain backgrounds. All images were resized to 224×224 pixels for consistency and efficiency. The core training set (5,000 images) focused on healthy and late blight classes from Kaggle. For testing, we used real-world field images from Tanzania (Mbeya, Arusha, Morogoro) to assess performance under natural conditions [20]. To improve diversity, we enriched the dataset with web-scraped images using tools like BeautifulSoup, Selenium, and requests. This hybrid approach boosted volume, class balance, and model generalization [21]. Table 1 shows dataset details; Figure 1 illustrates sample images.

The table 1 illustrates the dataset composition in detail, while Figure 1 presents examples of diseased and healthy leaves used for model training. All figures are explicitly referenced in the text in accordance with academic writing standards.




Figure 1: Examples of leaf diseases and healthy leaves

3.2 Data pre-processing

Data preprocessing is a crucial step for the efficient training of our CNN-based plant disease detection model, especially with a dataset of about 30 GB [22]. The images, acquired directly from the agricultural field to capture the variability of real-world conditions, were organized into a structured database, divided into training, validation, and test sets.

Several key strategies were applied during preprocessing:

- Resizing all images to a uniform size of 224 × 224 pixels to match the CNN input requirements.
- Normalization of pixel values to ensure faster convergence and numerical stability during training.
- Data augmentation techniques, including random rotations, horizontal flips, and zoom operations, were applied to artificially expand the dataset and improve the model's robustness to variations in orientation and scale.


In addition, binary label encoding was used to classify each image as either healthy (0) or diseased (1). This simplified the classification task and suited the structure of our binary CNN output layer. These preprocessing steps were instrumental in achieving excellent model performance and avoiding overfitting.

3.3 Model Architecture

Figure 2: System Architecture of the Fine-Tuned MobileNetV2 Pipeline for Early Plant Disease Detection

In this work, we introduce our end-to-end system architecture for early plant disease detection, centered on a fine-tuned MobileNetV2 model and its web-app integration. Figure 2 summarizes the complete workflow—from image acquisition and preprocessing through model inference to the user-facing output. The selection of an appropriate neural network architecture is a key element in building an effective plant disease detection system. In our work, we adopted MobileNetV2, a lightweight and efficient convolutional architecture known for its excellent compromise between accuracy and computational cost [23]. It is particularly suitable for deployment on resource-constrained platforms such as mobile or embedded devices.

We employed a transfer learning approach by using a version of MobileNetV2 pretrained on the



ImageNet dataset, serving as a fixed feature extractor. The original classification head was removed, and we added a custom classification layer adapted to our task. The final architecture is composed of the following layers:

- A MobileNetV2 base, with frozen weights to preserve its learned representations.
- A GlobalAveragePooling layer to reduce feature maps and prepare them for classification.
- A Dense layer for multi-class prediction.
- A Dropout layer to reduce overfitting.
- A final Dense layer producing raw logits for each class (softmax applied during loss computation).

This modified architecture enables the model to benefit from robust pre-learned features while adapting to our specific dataset of plant diseases.

The model was trained using the Adam optimizer, with a learning rate of 0.001, over 5 epochs and a batch size of 32. The input images were all resized to 224×224 pixels, normalized, and labeled using integer encoding, where each disease class was associated with a unique numerical label.

Once trained, the model demonstrated excellent generalization performance. On the test set, it achieved:

- Accuracy: 97.40%

- Loss: 0.1355

This performance confirms the effectiveness of the chosen architecture and training strategy, making it well-suited for practical applications in precision agriculture. The trained model and label mappings were saved to facilitate later reuse for prediction tasks on new unseen images.

3.4 Web Deployment

For the web deployment, we will develop a backend using a lightweight Python framework like Flask to create an API. The operational flow will begin when a user uploads an image through a web interface. This image is then sent to our backend, where it undergoes the same preprocessing steps used during training. The preprocessed image is fed into the loaded model to generate a prediction. To ensure a fast response time and efficient resource management, our model loading logic will involve loading the trained model into memory only once when the server starts, rather than on each individual request. Finally, the prediction output is returned to the user through the API, completing the cycle.

4 Results & Evaluation

This section details the experimentation process carried out to select the most appropriate architecture for the development of our plant disease detection web application. The main objective was to obtain a high-performance model in terms of accuracy while considering the constraints related to web deployment, particularly in terms of computational resources. Several convolutional network (CNN) architectures were evaluated, including AlexNet, a deep CNN architecture (DCNN), an augmented version of DCNN (DCNN(Aug)), ResNet, RCNN, and MobileNet. These architectures were chosen because of their recognized performance in similar image classification and detection tasks.

To implement our deep learning approach for plant disease detection, we relied on a cloud-based computational environment to ensure efficiency, flexibility, and scalability. The experimental setup is described below.

Tools and Environment:

All experiments were carried out using Python, taking advantage of its clarity, cross-platform compatibility, and rich ecosystem of AI libraries. For model development and training, we used:

- **TensorFlow/Keras:** for model definition, training, and evaluation.
- **Google Colaboratory (Colab):** as the main execution environment, providing free access to GPU acceleration.
- **Python libraries:** such as NumPy, Matplotlib, and Pandas for data manipulation and visualization.

The choice of Google Colab allowed us to handle large datasets (~30 GB) and train deep learning models without the need for powerful local hardware.

Training Configuration:

The model training was performed with the following settings:

- **Optimizer:** Adam
- **Loss function:** Binary Cross-Entropy for binary classification tasks, and Categorical Cross-Entropy for multi-class classification.
- **Input image size:** 224×224
- **Batch size:** 32
- **Number of epochs:** 5
- **Learning rate:** 0.001 (default for Adam)

All input images were normalized and resized before feeding into the network. Labels were encoded as integers, suitable for the SparseCategoricalCrossentropy loss function.

In addition, the number of images per disease is limited for each plant, as shown in the respective tables. To facilitate reuse of the model, weights obtained after training and the MobileNet v2 architecture, as well as matches between numeric labels and disease names (stored in a format), are saved in files and then compressed into an archive. This then loads the saved model and labels to make predictions on new image. The performance of the different architectures was compared using key metrics such as Accuracy and Loss, both on the training and validation sets. The results are presented in Table 4, the performance

on the test set is shown in Table 4:

Table 4: Model Performance on Training, Validation, and Test Sets.

Model	Training and Validation				Test Set		
	Train Acc.	Val Acc.	Train Loss	Val Loss	True Pred.	Accuracy	F1 Score
AlexNet	99.35%	92.05%	0.02	0.43	556	0.98	0.98
DCNN	100%	97.51%	3e-9	0.25	568	1.00	1.00
DCNN(Aug)	98.13%	96.53%	0.06	0.11	527	0.92	0.92
ResNet	92.25%	92.18%	1.57	1.57	517	0.91	0.91
RCNN	99.80%	93.55%	0.006	0.29	545	0.96	0.96
MobileNet	97.40%	97%	7e-6	0.03	570	0.97	0.97

Analysis of Tables 4 shows that MobileNet performs comparable to or better than other architectures tested, especially in terms of accuracy in validation and test data. Although DCNN and MobileNet achieve 100% accuracy on test data, MobileNet has the advantage of being much lighter in terms of the number of parameters (see example for Apple below), making it more suitable for deployment on a web app and mobile devices with limited resources. In addition, MobileNet achieves this performance with a reasonable number of epochs (20), which reduces training time.




Figure 3: Training and validation accuracy of the MobileNet model over epochs.

4.1 System Implementation

This section describes the practical implementation of the system, detailing the end-to-end user workflow, the backend architecture, the user interface, and the observed performance.

4.1.1 End-to-End Flow

The complete system operates on a simple and direct flow for the user: Input: The user selects and uploads an image of the plant in question through the web interface. Backend: Upon submission, an API call is made to the server. The backend, powered by our MobileNet model, receives the image, preprocesses it, and analyzes it to infer the probability of each disease. Output: The API returns a structured response containing the predicted disease label with the highest probability, along with the associated confidence score.

4.1.2 Interface Screenshots

The user interface was designed to be intuitive. The screenshots below illustrate the two main components of the user interaction: Web UI: Shows the main page where the user can upload their image. Response Box: Clearly displays the result returned by the model after analysis, including the disease name and the confidence score.

4.1.3 Average Runtime Latency

The system's performance was evaluated by measuring the total response time, which is the delay between the user submitting the image and receiving the prediction. Average runtime latency: The average response time observed is 0.2 seconds.




Figure a: Interface for selecting the plant type. **Figure b:** Interface for choosing a specific disease to view information. **Figure c:** Image gallery of plant diseases.

5 Conclusion and Future Work

5.1 Conclusion

This project successfully addressed the challenge of creating a rapid and accessible tool for plant disease identification. The primary achievement is the development and validation of a lightweight classifier built upon the MobileNet architecture. This choice proved highly effective, striking an optimal balance between computational efficiency and predictive power. The resulting model not only achieved high classification accuracy but did so with a small footprint, making it ideal for deployment in resource-constrained environments.

5.2 Future Work

While the current system provides a robust proof-of-concept, several key enhancements are planned to expand its capabilities and real-world impact. The following areas represent our primary focus for future development: Expansion of the Disease Database:

The model's current utility is limited to the diseases included in its training set. A critical next step is to significantly expand the dataset to include a wider variety of diseases (e.g., bacterial spot, rusts, and various blights) across a more diverse range of plant species. This will require substantial data collection and annotation efforts but will drastically improve the system's robustness and practical value in diverse agricultural settings.

References

- [1] FAO, "Plant health and food security," 2021.
- [2] S. Sankaran et al., *Comput. Electron. Agric.*, 72(1):1–13, 2010.
- [3] Y. Xie et al., *Mol. Plant*, 8(8):1274–1284, 2015.
- [4] R. Benard et al., *J. Inf. Sci. Technol.*, 4:1–11, 2014.
- [5] A. Kamaras & F. Prenafeta-Boldú, *Comput. Electron. Agric.*, 147:70–90, 2018.
- [6] Y. Han et al., *Environ. Res.*, 208:112761, 2021.
- [7] K. Ferentinos, *Comput. Electron. Agric.*, 145:311–318, 2018.
- [8] A. G. Howard et al., arXiv:1704.04861, 2017.
- [9] M. Sandler et al., CVPR:4510–4520, 2018.
- [10] A. H. Kulkarni, IJCSMC, 5(6):1–7, 2016.
- [11] S. Sladojevic et al., *Comput. Intell. Neurosci.*, 2016:1–11.
- [12] D. P. Hughes & M. Salathé, arXiv:1511.08060, 2015.
- [13] S. P. Mohanty et al., *Front. Plant Sci.*, 7:1419, 2016.
- [14] R. A. Mique Jr. & R. C. Eusebio, IJEAT, 8(5):846–850, 2019.
- [15] A. Ramcharan et al., *Front. Plant Sci.*, 10:272, 2019.
- [16] A. Rajbongshi et al., ISMSIT:1–7, 2020.
- [17] E. H. El Fatimi et al., IEEE Access, 2022.

- [18] Q. Zhu et al., AAAI:5973–5980, 2019.
- [19] S. Ashwinkumar et al., Mater. Today Proc., 51:1–7, 2022.
- [20] J. S. Fuentes, M. De Bei, M. Pech, and S. Tyerman, “Automated grapevine berry size prediction using image analysis,” *Sensors*, vol. 14, no. 12, pp. 22215–22235, 2014.
- [21] J. Rumpf, A. Mahlein, U. Steiner, E.-C. Oerke, H.-W. Dehne, and L. Plümer, “Early detection and classification of plant diseases with support vector machines based on hyperspectral reflectance,” *Computers and Electronics in Agriculture*, vol. 74, no. 1, pp. 91–99, 2010.
- [22] A. A. Barbedo, “A review on the main challenges in automatic plant disease identification based on visible range images,” *Biosystems Engineering*, vol. 144, pp. 52–60, 2016.
- [23] M. Brahimi, K. Boukhalfa, and A. Moussaoui, “Deep learning for tomato diseases: Classification and symptoms visualization,” *Applied Artificial Intelligence*, vol. 31, no. 4, pp. 299–315, 2017.
- [24] M. Jeong, Y. Lee, and Y. Hwang, “Plant disease detection and classification based on RGB and hyperspectral imaging technologies,” *Journal of Sensors*, vol. 2019, Art. no. 6358145, pp. 1–15, 2019.